



universidade de aveiro
theoria poiesis praxis

Camada Protocolar de Aplicação (Application Layer)

Redes de Comunicações 1

Licenciatura em Engenharia de Computadores e
Informática

DETI-UA, 2023/2024

TFTP - TRIVIAL FILE TRANSFER PROTOCOL

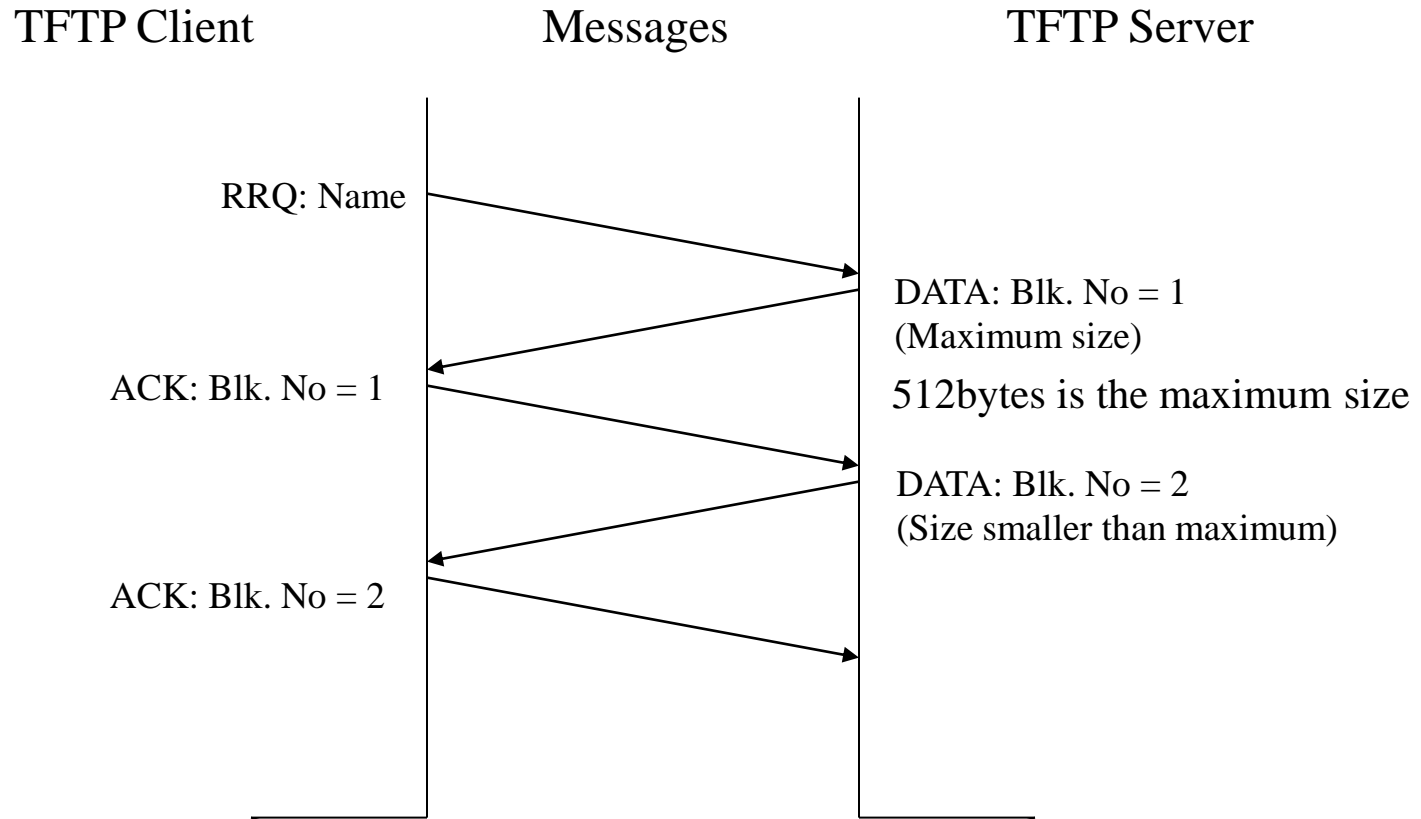
Trivial File Transfer Protocol (TFTP)

- Transfer file service with very simple client/server interactions.
- A TFTP client, to receive a file, needs to know the name and directory where the file is stored.
- This service can be used to configure network elements.
- Most *Routers* and *Switches* allow its configuration to be performed through a TFTP server to receive the configuration file.
 - They need the IP address of the server and the name of the file.
 - Very small files.

Trivial File Transfer Protocol (TFTP)

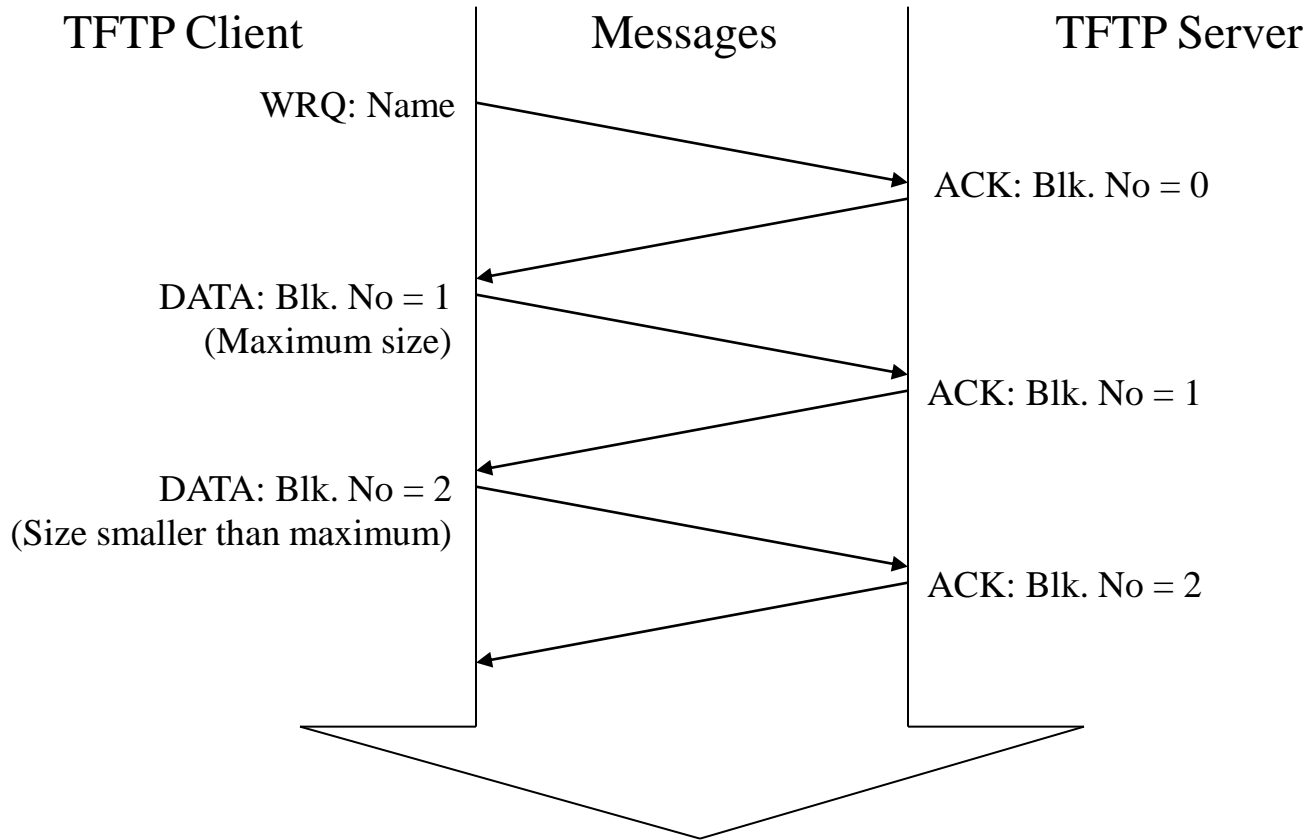
- Basic file transfer service (IETF RFC 1350)
 - It does not allow to list server files
 - It does not support authentication
- TFTP runs over **UDP**
 - The initial client message is sent to server port number **69**
 - **TFTP is the one that has to address packet retransmission**
 - TFTP server answers from another locally selected port number
 - The following messages are exchanged with the selected port number
- TFTP uses *Stop and Wait* ARQ mechanism
- TFTP has 5 messages:
 - *Read Request (RRQ)*
 - *Write Request (WRQ)*
 - *Data*
 - *Acknowledgement (ACK)*
 - *Error (ERR)*

Read Request session



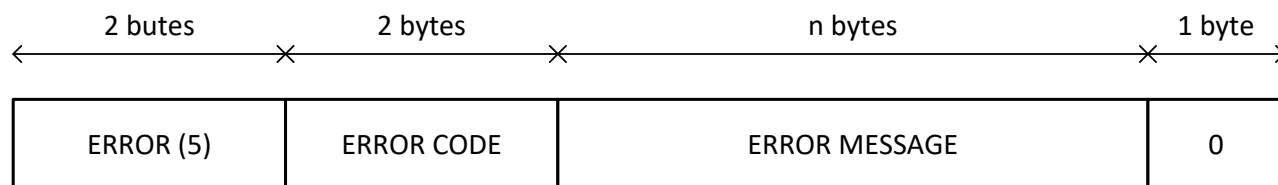
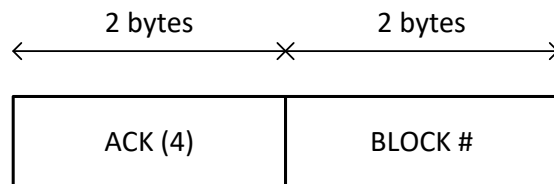
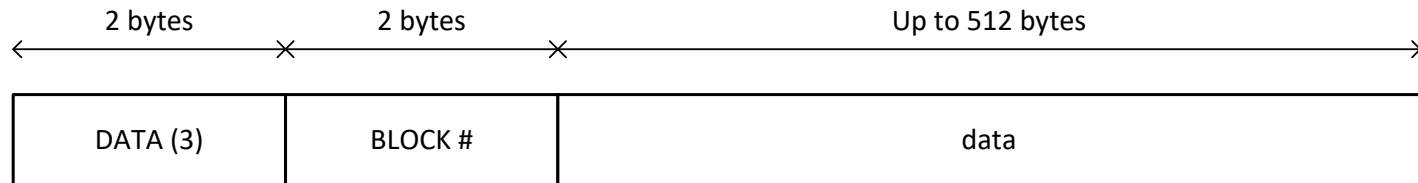
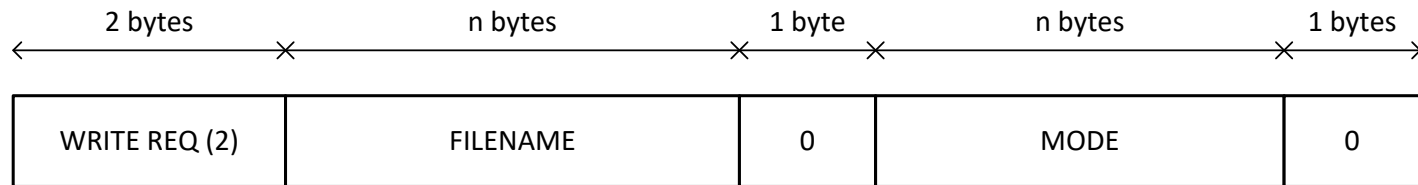
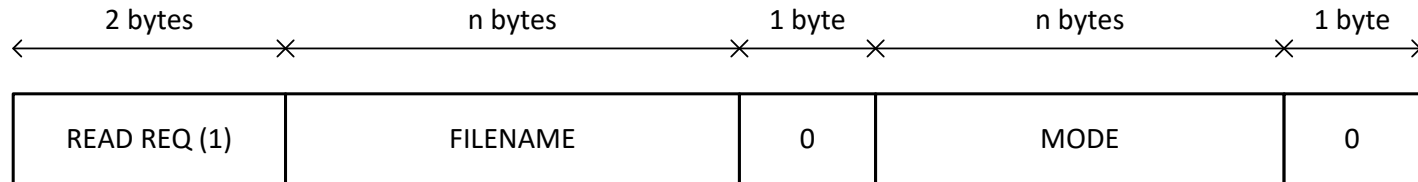
Client detects the last block of DATA when the data block size is lower than the maximum size. If the size of the file is **multiple** of the maximum size, an additional data block is sent with **0** bytes of data. RRQ always with odd number of packets (RRQ+pairs of data and ACK).

Write Request session




Server detects the last block of DATA when the data block size is lower than the maximum size. If the size of the file is **multiple** of the maximum size, an additional data block is sent with **0** bytes of data. WRQ always with even number of packets (WRQ+ACK0+pairs of data and ACK).

Messages format



FTP – FILE TRANSFER PROTOCOL


FTP: the file transfer protocol



- ❑ transfer file to/from remote host (IETF RFC 959)
- ❑ client/server model
 - ❖ *client*: side that initiates file transfer (either to/from remote host)
 - ❖ *server*: remote host
- ❑ ftp: runs over **TCP** (TCP takes care of packet retransmission)
- ❑ ftp server ports: **21** and **20**
 - ❖ **21** for the control connection
 - ❖ **20** for each data connection

FTP: separate control, data connections

- ❑ FTP client contacts FTP server at port 21, TCP is transport protocol
- ❑ client authorized over control connection
- ❑ client browses remote directory by sending commands over control connection
- ❑ when server receives file transfer command, server opens 2nd TCP connection (for file transfer) to client
- ❑ after transferring one file, server closes data connection.



- ❑ server opens another TCP data connection to transfer another file.
- ❑ control connection: "out of band"
- ❑ FTP server maintains the "state" of client interactions: current directory, earlier authentication, etc...

FTP commands, responses

Sample commands:

- ❑ sent as ASCII text over control channel:
- ❑ `USER username`
- ❑ `PASS password`
- ❑ `LIST` (return list of files in current directory)
- ❑ `RETR filename` (retrieves file from server)
- ❑ `STOR filename` (stores file onto server)

Sample return codes

- ❑ status code and phrase (as in HTTP):
- ❑ 331 Username OK, password required
- ❑ 125 Data connection already open; transfer starting
- ❑ 425 Can't open data connection
- ❑ 452 Error writing file

FTP client: user commands

```
MS Command Prompt - ftp
Microsoft(R) Windows NT(TM)
(C) Copyright 1985-1996 Microsoft Corp.

C:\>ftp
ftp> help
Commands may be abbreviated.  Commands are:

!                delete          literal          prompt          send
?                debug           ls              put             status
append          dir             mdelete        pwd            trace
ascii          disconnect    mdir           quit           type
bell           get           mget          quote         user
binary         glob          mkdir         recv         verbose
bye            hash          mls           remotehelp
cd             help          mput          rename
close          lcd           open          rmdir
ftp> _
```

Initial connection to FTP server

The screenshot displays two windows from a Windows NT environment. The top window is NetXray, showing a packet capture on the interface 'Local/3Com EtherLink XL COMBO 10Mb Ethernet NIC (3C900-COMBO)_1'. The packet list shows five packets related to an FTP connection. The bottom window is a Command Prompt titled 'Command Prompt - ftp 192.168.8.235', showing the successful execution of the 'ftp' command to connect to the server at 192.168.8.235.

No.	Sta.	Source Address	Dest Address	Layer	Len	Summary
1	Ok	192.168.8.227	192.168.8.235	TCP	62	1246->File Transfer (Control),[Syn]
2	Ok	192.168.8.235	192.168.8.227	TCP	64	File Transfer (Control)->1246,[Syn]
3	Ok	192.168.8.227	192.168.8.235	TCP	58	1246->File Transfer (Control),S=152'
4	Ok	192.168.8.235	192.168.8.227	FTP	111	220 labcomSIII Microsoft FTP Service
5	Ok	192.168.8.227	192.168.8.235	TCP	58	1246->File Transfer (Control),S=152'

Transmission Control Protocol
Port 1246 ---> File Transfer (Control)

```
Microsoft(R) Windows NT(TM)
(C) Copyright 1985-1996 Microsoft Corp.

C:\>ftp 192.168.8.235
Connected to 192.168.8.235.
220 labcomSIII Microsoft FTP Service (Version 5.0).
User (192.168.8.235:(none)): _
```

TCP 3-way handshake to initiate the control connection in port 21

For each packet sent, a TCP ACK is needed to confirm the previous packet

Introduction of username

The screenshot displays two windows from a Windows NT environment. The top window, titled "NetXRay - Local/3Com EtherLink XL COMBO 10Mb Ethernet NIC [3C900-COMBO]_1 - [XRay2 : 1/...", contains a packet capture table and a protocol tree. The table lists three packets: a USER command, a password request, and a file transfer control packet. The protocol tree shows the Transmission Control Protocol and its connection to Port 1246. The bottom window is a Command Prompt titled "Command Prompt - ftp 192.168.8.235", showing the execution of the ftp command and the resulting FTP session output, including the password prompt.

No.	Sta.	Source Address	Dest Address	Layer	Len	Summary
1	Ok	192.168.8.227	192.168.8.235	FTP	74	USER anonymous
2	Ok	192.168.8.235	192.168.8.227	FTP	96	331 Password required for anonymous
3	Ok	192.168.8.227	192.168.8.235	TCP	58	1246->File Transfer (Control).S=152'

Transmission Control Protocol
Port 1246 ---> File Transfer (Control)

```
Microsoft(R) Windows NT(TM)  
(C) Copyright 1985-1996 Microsoft Corp.  
  
C:\>ftp 192.168.8.235  
Connected to 192.168.8.235.  
220 labcomSIII Microsoft FTP Service (Version 5.0).  
User (192.168.8.235:(none)): anonymous  
331 Password required for anonymous.  
Password: _
```

Introduction of password

The screenshot displays two windows from a Windows NT environment. The top window is NetXRay, showing a packet capture on the interface 'Local/3Com EtherLink XL COMBO 10Mb Ethernet NIC [3C900-COMBO]_1'. The packet list shows three packets: a PASS command, a 230 login response, and a 1249 control connection. The bottom window is a Command Prompt titled 'Command Prompt - ftp 192.168.8.235', showing the execution of the 'ftp' command and the resulting FTP session output.

NetXRay - Local/3Com EtherLink XL COMBO 10Mb Ethernet NIC [3C900-COMBO]_1 - [XRay3 : 1/...

File Capture Packet Tools Window Help

No.	Sta.	Source Address	Dest Address	Layer	Len	Summary
1	Ok	192.168.8.227	192.168.8.235	FTP	76	PASS curso@ua.pt
2	Ok	192.168.8.235	192.168.8.227	FTP	89	230 Anonymous user logged in.
3	Ok	192.168.8.227	192.168.8.235	TCP	58	1249->File Transfer (Control).S=153'

Transmission Control Protocol

Port 1249 ---> File Transfer (Control)

Command Prompt - ftp 192.168.8.235

```
Microsoft(R) Windows NT(TM)
(C) Copyright 1985-1996 Microsoft Corp.

C:\>ftp 192.168.8.235
Connected to 192.168.8.235.
220 labcomSIII Microsoft FTP Service (Version 5.0).
User (192.168.8.235:(none)): anonymous
331 Anonymous access allowed, send identity (e-mail name) as password.
Password:
230 Anonymous user logged in.
ftp> _
```

Introduction of DIR command

No.	Sta.	Source Address	Dest Address	Layer	Len	Summary
1	Ok	192.168.8.227	192.168.8.235	FTP	84	PORT 192.168.8.227,4,229
2	Ok	192.168.8.235	192.168.8.227	FTP	88	200 PORT command successful.
3	Ok	192.168.8.227	192.168.8.235	FTP	64	LIST
4	Ok	192.168.8.235	192.168.8.227	FTP	111	150 Opening ASCII mode data connect
5	Ok	192.168.8.235	192.168.8.227	TCP	66	File Transfer (Default Data)->1253,
6	Ok	192.168.8.227	192.168.8.235	TCP	62	1253->File Transfer (Default Data),
7	Ok	192.168.8.235	192.168.8.227	TCP	64	File Transfer (Default Data)->1253,
8	Ok	192.168.8.235	192.168.8.227	FTP	210	Data (total 152 bytes), (More data)
9	Ok	192.168.8.235	192.168.8.227	TCP	64	File Transfer (Default Data)->1253,
10	Ok	192.168.8.227	192.168.8.235	TCP	58	1253->File Transfer (Default Data),
11	Ok	192.168.8.227	192.168.8.235	TCP	58	1253->File Transfer (Default Data),
12	Ok	192.168.8.235	192.168.8.227	TCP	64	File Transfer (Default Data)->1253,
13	Ok	192.168.8.227	192.168.8.235	TCP	58	1252->File Transfer (Control),S=154,
14	Ok	192.168.8.235	192.168.8.227	FTP	82	226 Transfer complete.
15	Ok	192.168.8.227	192.168.8.235	TCP	58	1252->File Transfer (Control),S=154,

```

File Transfer Protocol
PORT 192.168.8.227,4,229

MS-DOS Command Prompt - ftp 192.168.8.235
Microsoft(R) Windows NT(TM)
(C) Copyright 1985-1996 Microsoft Corp.

C:\>ftp 192.168.8.235
Connected to 192.168.8.235.
220 labcomSIII Microsoft FTP Service (Version 5.0).
User (192.168.8.235:(none)): anonymous
331 Anonymous access allowed, send identity (e-mail name) as password.
Password:
230 Anonymous user logged in.
ftp> dir
200 PORT command successful.
150 Opening ASCII mode data connection for /bin/ls.
10-05-00 12:02PM <DIR> CursorTCP_IP
08-27-00 12:32AM 1024 f1024.txt
08-27-00 12:32AM 1268 f1268.txt
226 Transfer complete.
152 bytes received in 0.00 seconds (152000.00 Kbytes/sec)
ftp>
  
```

Information to start the data connection

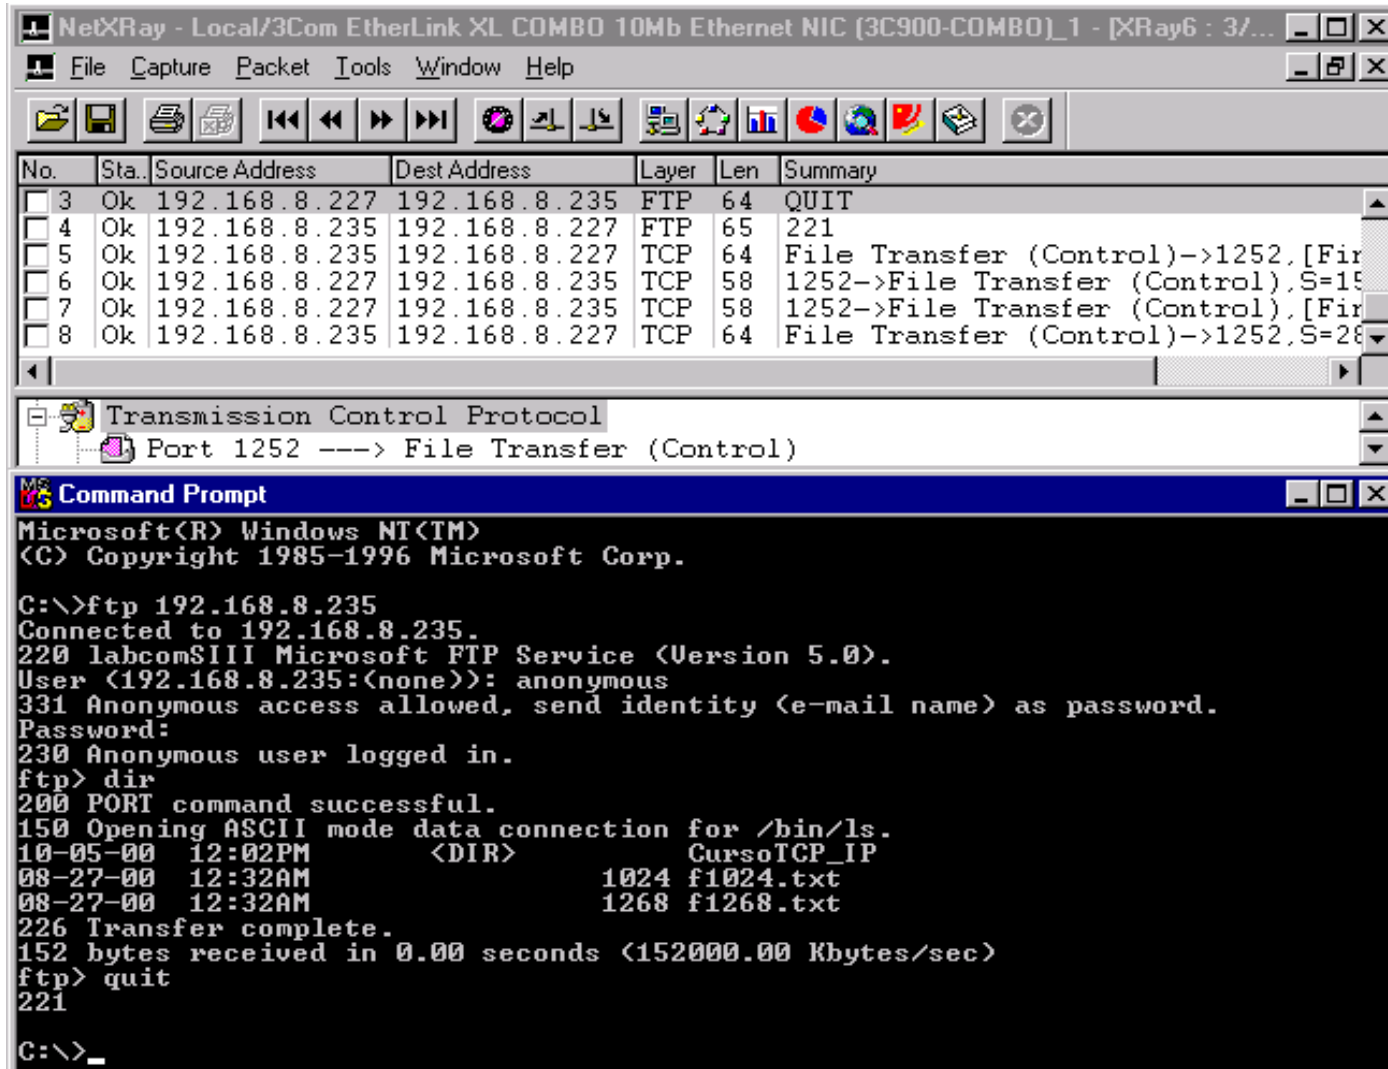
FTP command for DIR

TCP 3-way handshake to initiate the data connection in port 20

Data is sent

FIN and ACK packets (4 in total) to finish the data connection in port 20

Termination of connection (quit)



The screenshot displays two windows. The top window, 'NetXray - Local/3Com EtherLink XL COMBO 10Mb Ethernet NIC [3C900-COMBO]_1 - [XRay6 : 3/...', shows a packet capture table. The bottom window, 'Command Prompt', shows the execution of an FTP session.

No.	Sta.	Source Address	Dest Address	Layer	Len	Summary
3	Ok	192.168.8.227	192.168.8.235	FTP	64	QUIT
4	Ok	192.168.8.235	192.168.8.227	FTP	65	221
5	Ok	192.168.8.235	192.168.8.227	TCP	64	File Transfer (Control)->1252,[Fir
6	Ok	192.168.8.227	192.168.8.235	TCP	58	1252->File Transfer (Control),S=15
7	Ok	192.168.8.227	192.168.8.235	TCP	58	1252->File Transfer (Control),[Fir
8	Ok	192.168.8.235	192.168.8.227	TCP	64	File Transfer (Control)->1252,S=28

Below the table, the 'Transmission Control Protocol' section shows 'Port 1252 ----> File Transfer (Control)'.

```
Microsoft(R) Windows NT(TM)
(C) Copyright 1985-1996 Microsoft Corp.

C:\>ftp 192.168.8.235
Connected to 192.168.8.235.
220 labcomSIll Microsoft FTP Service (Version 5.0).
User (192.168.8.235:(none)): anonymous
331 Anonymous access allowed, send identity (e-mail name) as password.
Password:
230 Anonymous user logged in.
ftp> dir
200 PORT command successful.
150 Opening ASCII mode data connection for /bin/ls.
10-05-00 12:02PM <DIR> CursorTCP_IP
08-27-00 12:32AM 1024 f1024.txt
08-27-00 12:32AM 1268 f1268.txt
226 Transfer complete.
152 bytes received in 0.00 seconds (152000.00 Kbytes/sec)
ftp> quit
221
C:\>_
```

FIN and ACK packets (4 in total) to finish the control connection in port 21

DNS – DOMAIN NAME SYSTEM

DNS: Domain Name System

People: many identifiers:

- ❖ SSN, name, passport #

Internet hosts, routers:

- ❖ IP address (32 bit) - used for addressing datagrams
- ❖ "name" (for example, www.yahoo.com) - used by humans


DNS: provides a mapping between IP addresses and names

Works in **UDP port 53**

Domain Name System:

- ❑ *distributed database*
implemented in a hierarchy of many *Name Servers*
- ❑ *application-layer protocol*
hosts communicate with Name Servers to *resolve* names (name - IP address translation)
 - ❖ note: core Internet function, implemented as application-layer protocol
 - ❖ complexity at network's "edge"

Distributed, Hierarchical Database



Client wants IP for **www.amazon.com**; 1st approximation:

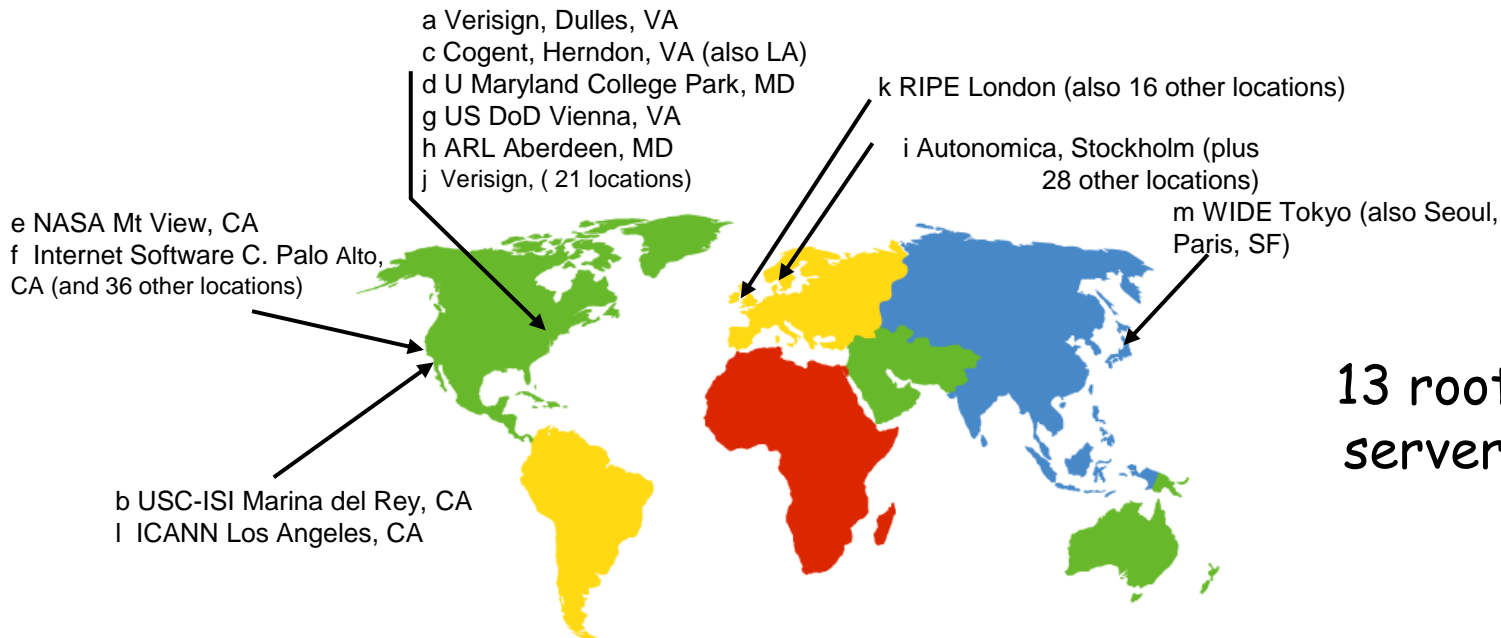
- ❑ client queries a root server to find 'com' DNS server
- ❑ client queries 'com' DNS server to get 'amazon.com' DNS server
- ❑ client queries 'amazon.com' DNS server to get IP address for **www.amazon.com**

DNS: Root Name Servers

Server	Operator	Cities	IP Addr	Home AS	Answers ICMP?
A	VeriSign Global Registry Services	Herndon VA, US	198.41.0.4	19836	<i>yes</i>
B	Information Sciences Institute	Marina Del Rey CA, US	128.9.0.107	<i>tba</i>	<i>yes</i>
C	Cogent Communications	Herndon VA, US	192.33.4.12	2149	<i>yes</i>
D	University of Maryland	College Park MD, US	128.8.10.90	27	<i>yes</i>
E	NASA Ames Research Center	Mountain View CA, US	192.203.230.10	297	<i>yes</i>
F	Internet Software Consortium	Palo Alto CA, US; San Francisco CA, US; Madrid, ES; San Jose, CA, US; New York, NY, US; Hong Kong, HK	IPv4: 192.5.5.241 IPv6: 2001:500::1035	3557	<i>yes</i>
G	U.S. DOD Network Information Center	Vienna VA, US	192.112.36.4	568	<i>no</i>
H	U.S. Army Research Lab	Aberdeen MD, US	IPv4: 128.63.2.53 IPv6: 2001:500:1::803f:235	13	<i>yes</i>
I	Autonomica	Stockholm, SE	192.36.148.17	8674	<i>yes</i>
J	VeriSign Global Registry Services	Herndon VA, US	192.58.128.30	26415	<i>yes</i>
K	Reseaux IP Europeens - Network Coordination Centre	London, UK	193.0.14.129	5459	<i>yes</i>
L	Internet Corporation for Assigned Names and Numbers	Los Angeles CA, US	198.32.64.12	20144	<i>no</i>
M	WIDE Project	Tokyo, JP	202.12.27.33	7500	<i>yes</i>

DNS: Root Name Servers

- contacted by local Name Server that cannot resolve name



13 root name
servers worldwide

Top Level Domains (TLD)

❑ Organizational domains:

- ❖ com - commercial organizations
- ❖ edu - educational institutions
- ❖ gov - govern institutions
- ❖ mil - military institutions
- ❖ net - network operators
- ❖ int - international organizations
- ❖ org - other organizations

❑ Country domains:

- ❖ pt - Portugal
- ❖ es - Spain

Local Name Server

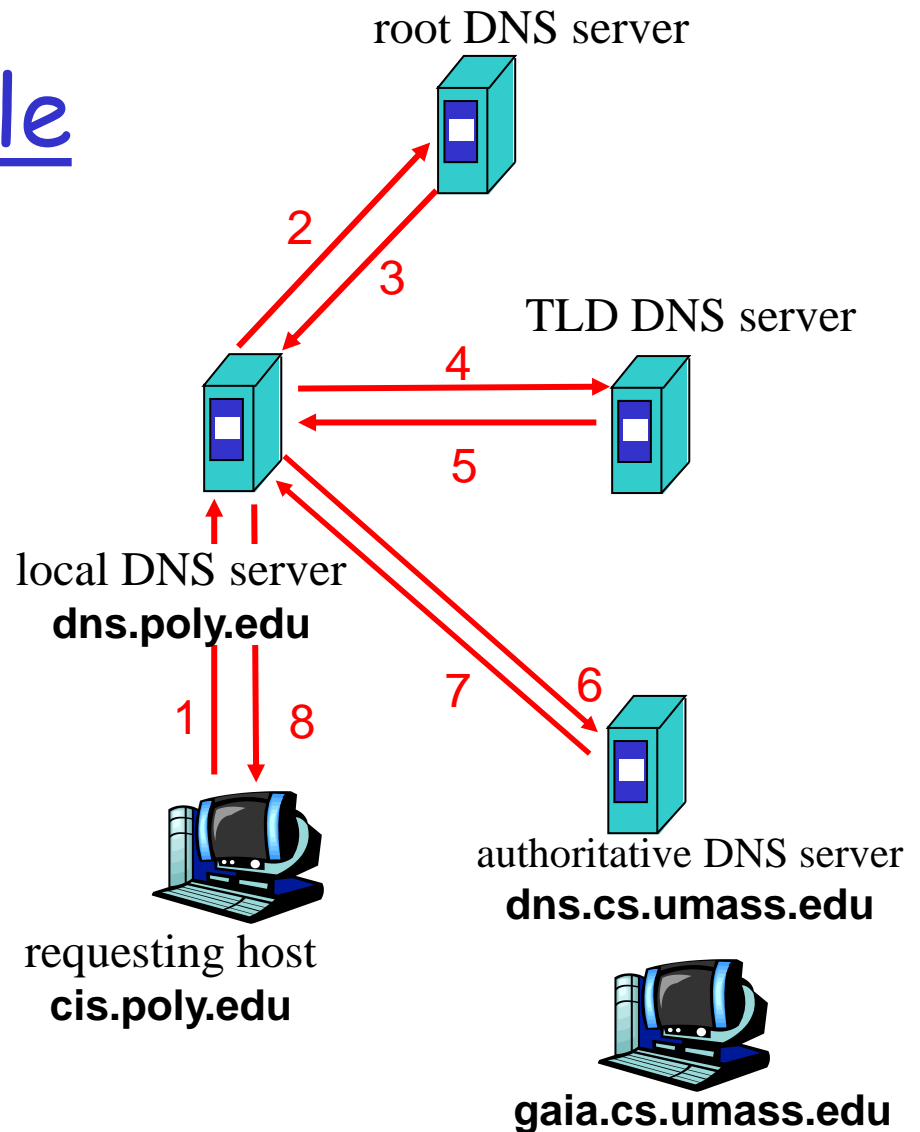
- ❑ does not strictly belong to hierarchy
- ❑ each ISP (residential ISP, company, university) has one.
 - ❖ also called "default Name Server"
- ❑ when host makes DNS query, query is sent to its local DNS server
 - ❖ acts as proxy, forwards query into hierarchy

DNS name resolution example

- Host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:

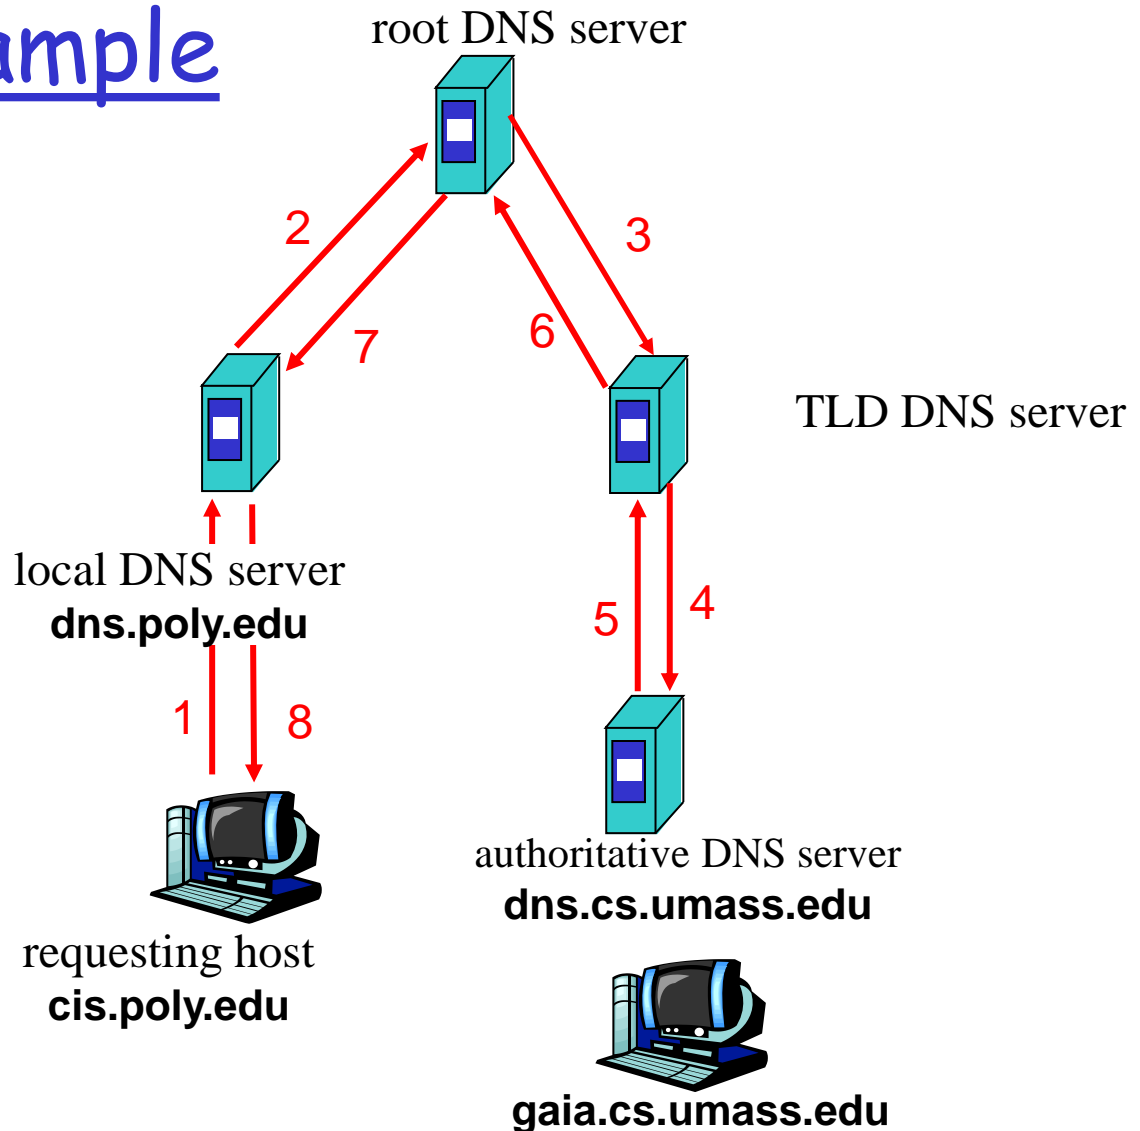
- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"



DNS name resolution example

recursive query:

- ❑ puts burden of name resolution on contacted name server



Recursive vs. iterated resolution

❑ Recursive resolution:

- ❖ More efficient: minimizes the time between the query and the answer
- ❖ Requires more processing power in DNS servers: each server has more simultaneous ongoing requests, on average

❑ Iterated resolution:

- ❖ Less efficient: the time between the query and the answer is larger, on average
- ❖ Minimizes the processing power required on DNS servers: each server replies immediately to each received query

DNS: caching and updating records

- once (any) Name Server learns mapping, it *caches* mapping
 - ❖ cache entries timeout (disappear) after some time
 - ❖ TLD server addresses are typically cached in local Name Servers
 - Thus, Root Name Servers are not often visited

DNS records

DNS: distributed database storing Resource Records (RR)

RR format: (name, value, type, ttl)

□ Type=A

- ❖ **name** is hostname
- ❖ **value** is IP address
- ❖ e.g. (relay1.bar.foo.com, 145.37.93.126, A)

□ Type=NS

- ❖ **name** is domain (e.g. foo.com)
- ❖ **value** is hostname of authoritative name server for this domain
- ❖ e.g. (foo.com, dns.foo.com, NS)

□ Type=CNAME

- ❖ **name** is alias name for some "canonical" (the real) name
- ❖ **value** is canonical name
- ❖ e.g. (foo.com, relay1.bar.foo.com, CNAME)

□ Type=MX

- ❖ **value** is name of mailserver associated with **name**
- ❖ e.g. (foo.com, mail.bar.foo.com, MX)

□ Type=AAAA

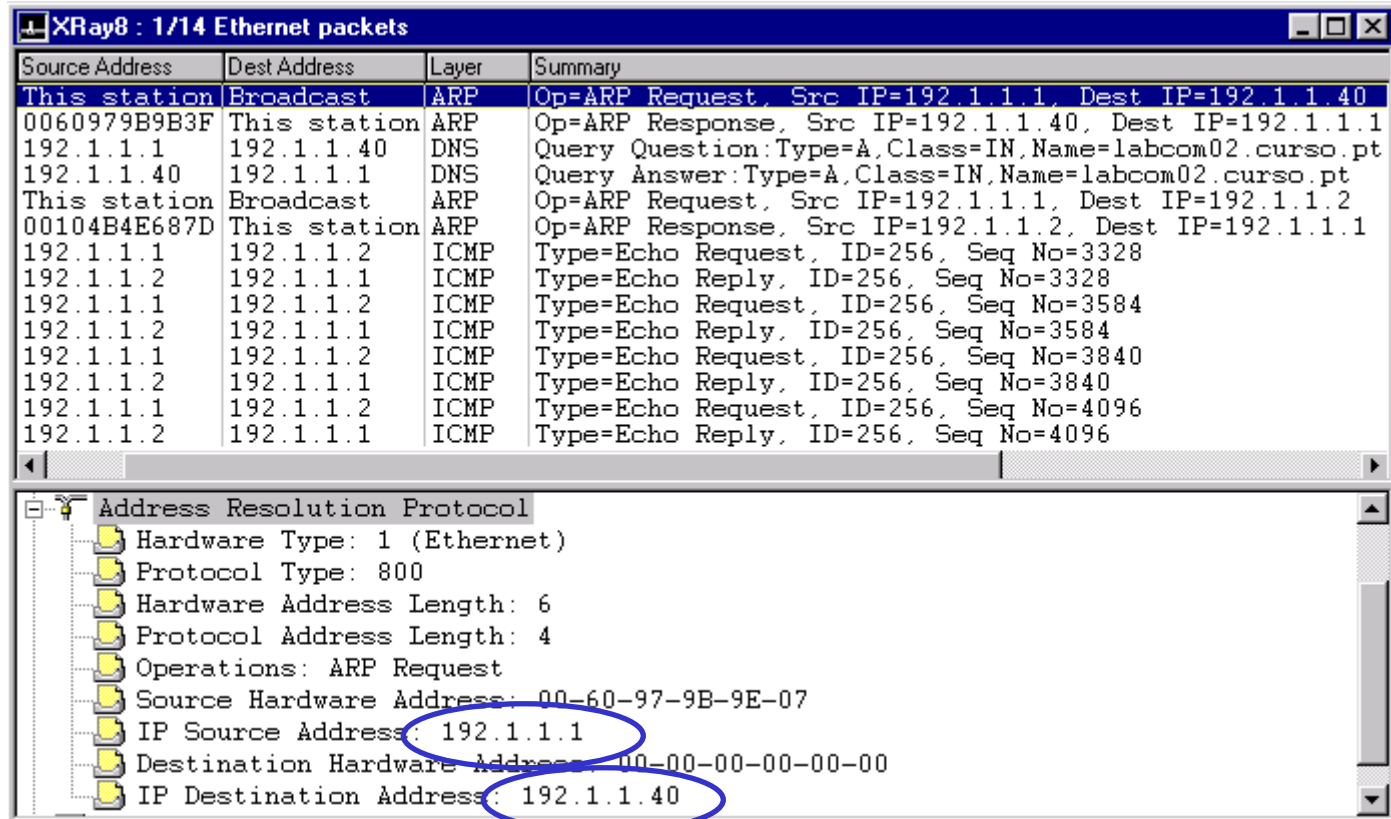
- ❖ **name** is hostname
- ❖ **value** is IPv6 address

DNS messages

host: 192.1.1.1

DNS: 192.1.1.40

C:\>ping labcom02.curso.pt



The screenshot shows the XRay8 network analysis tool. The top pane displays a list of 14 Ethernet packets. The bottom pane shows the details for the selected packet (packet 1, ARP Request), with the IP Source Address (192.1.1.1) and IP Destination Address (192.1.1.40) circled in blue.

Source Address	Dest Address	Layer	Summary
This station	Broadcast	ARP	Op=ARP Request, Src IP=192.1.1.1, Dest IP=192.1.1.40
0060979B9B3F	This station	ARP	Op=ARP Response, Src IP=192.1.1.40, Dest IP=192.1.1.1
192.1.1.1	192.1.1.40	DNS	Query Question:Type=A,Class=IN,Name=labcom02.curso.pt
192.1.1.40	192.1.1.1	DNS	Query Answer:Type=A,Class=IN,Name=labcom02.curso.pt
This station	Broadcast	ARP	Op=ARP Request, Src IP=192.1.1.1, Dest IP=192.1.1.2
00104B4E687D	This station	ARP	Op=ARP Response, Src IP=192.1.1.2, Dest IP=192.1.1.1
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3328
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3328
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3584
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3584
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3840
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3840
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=4096
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=4096

Address Resolution Protocol	
Hardware Type:	1 (Ethernet)
Protocol Type:	800
Hardware Address Length:	6
Protocol Address Length:	4
Operations:	ARP Request
Source Hardware Address:	00-60-97-9B-9E-07
IP Source Address:	192.1.1.1
Destination Hardware Address:	00-00-00-00-00-00
IP Destination Address:	192.1.1.40

DNS messages

The screenshot shows the XRay8 network analysis tool. The top pane displays a list of 14 Ethernet packets. The bottom pane shows a detailed view of the selected packet (packet 3), which is a DNS query. The packet details include: User Datagram Protocol, Port 1052 (circled) to Domain Name Server, Total length: 43 (Octets), Checksum: 0x5446, Domain Name Service, HEADER SECTION, Identifier: 1, Flags: Req, Query, Non-Auth, Recu Desr, Recu Not Ava, RCode=No Error, Section Entries: QDCount=1, ANCount=0, NSCount=0, ARCount=0, and QUESTION SECTION[1]: Type=A, Class=IN, labcom02.cursor.pt (circled). The packet data is shown in hexadecimal and ASCII at the bottom.

Source Address	Dest Address	Layer	Summary
This station	Broadcast	ARP	Op=ARP Request, Src IP=192.1.1.1, Dest IP=192.1.1.40
0060979B9B3F	This station	ARP	Op=ARP Response, Src IP=192.1.1.40, Dest IP=192.1.1.1
192.1.1.1	192.1.1.40	DNS	Query Question:Type=A,Class=IN,Name=labcom02.cursor.pt
192.1.1.40	192.1.1.1	DNS	Query Answer:Type=A,Class=IN,Name=labcom02.cursor.pt
This station	Broadcast	ARP	Op=ARP Request, Src IP=192.1.1.1, Dest IP=192.1.1.2
00104B4E687D	This station	ARP	Op=ARP Response, Src IP=192.1.1.2, Dest IP=192.1.1.1
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3328
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3328
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3584
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3584
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3840
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3840
192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=4096
192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=4096

User Datagram Protocol

- Port 1052 --> Domain Name Server
- Total length: 43 (Octets)
- Checksum: 0x5446
- Domain Name Service
 - HEADER SECTION:
 - Identifier: 1
 - Flags: Req, Query, Non-Auth, Recu Desr, Recu Not Ava, RCode=No Error
 - Section Entries: QDCount=1, ANCount=0, NSCount=0, ARCount=0
 - QUESTION SECTION[1]: Type=A, Class=IN, labcom02.cursor.pt

00000000: 00 60 97 9b 9b 3f 00 60 97 9b 9e 07 08 00 45 00 |?.....E.

00000010: 00 3f bf 00 80 80 80 11 f9 81 c0 01 01 01 c0 01 | ?&.....A..A.

00000020: 01 28 04 7c 00 35 00 2b 54 46 00 01 01 00 00 01 | (...).5.+TF.....

00000030: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |

Source port: 1052

Destination port: 53

DNS messages

XRay8.cap : 4/14 Ethernet packets					
No.	Stat...	Source Address	Dest Address	Layer	Summary
1	Ok	This station	Broadcast	ARP	Op=ARP Request, Src IP=192.1.1.1, Dest IP=192
2	Ok	0060979B9B3F	This station	ARP	Op=ARP Response, Src IP=192.1.1.40, Dest IP=1
3	Ok	192.1.1.1	192.1.1.40	DNS	Query Question:Type=A,Class=IN,Name=labcom02.
4	Ok	192.1.1.40	192.1.1.1	DNS	Query Answer:Type=A,Class=IN,Name=labcom02.cu
5	Ok	This station	Broadcast	ARP	Op=ARP Request, Src IP=192.1.1.1, Dest IP=192
6	Ok	00104B4E687D	This station	ARP	Op=ARP Response, Src IP=192.1.1.2, Dest IP=19
7	Ok	192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3328
8	Ok	192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3328
9	Ok	192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3584
10	Ok	192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3584
11	Ok	192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=3840
12	Ok	192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=3840
13	Ok	192.1.1.1	192.1.1.2	ICMP	Type=Echo Request, ID=256, Seq No=4096
14	Ok	192.1.1.2	192.1.1.1	ICMP	Type=Echo Reply, ID=256, Seq No=4096

[-] User Datagram Protocol
[-] Port Domain Name Server ---> 1052
[-] Total length: 59 (Octets)
[-] Checksum: 0xA915
[-] Domain Name Service
[-] HEADER SECTION:
[-] Identifier: 1
[+] Flags: Resp, Query, Auth, Recu Desr, Recu Ava, RCode=No Error
[+] Section Entries: QDCount=1 , ANCount=1 , NSCount=0 , ARCount=0
[+] QUESTION SECTION[1]: Type=A, Class=IN, labcom02.cursor.pt
[+] ANSWER SECTION[1]: Type=A, Class=IN, 192.1.1.2

00000020:	01 01 00 35 04 1c 00 3b a9 15 00 01 85 80 00 01	...5.....
00000030:	00 01 00 00 00 00 08 6c 61 62 63 6f 6d 30 32 05labcom02.

HTTP

Web and HTTP

First some basic concepts

- ❑ Web page consists of objects
- ❑ Object can be HTML file, JPEG image, Java applet, audio file,...
- ❑ Web page consists of base HTML-file which includes several referenced objects
- ❑ Each object is addressable by a URL
- ❑ Example URL:

`www.someschool.edu/someDept/pic.gif`


host name

path name

HTTP overview

HTTP: HyperText Transfer Protocol

- Web's application layer protocol
- client/server model
 - ❖ *client*: browser that requests, receives, "displays" Web objects
 - ❖ *server*: Web server sends objects in response to requests
- HTTP 1.0: RFC 1945
- HTTP 1.1: RFC 2068



HTTP overview (continued)

Uses TCP:

- ❑ client initiates TCP connection (creates socket) to server, port 80
- ❑ server accepts TCP connection from client
- ❑ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ❑ TCP connection closed

HTTP is "stateless"

- ❑ server maintains no information about past client requests

aside
Protocols that maintain "state" are complex!

- ❑ past history (state) must be maintained
- ❑ if server/client crashes, their views of "state" may be inconsistent, must be reconciled

HTTP connections

Nonpersistent HTTP

- ❑ At most one object is sent over a TCP connection
- ❑ HTTP/1.0 uses nonpersistent HTTP

Persistent HTTP

- ❑ Multiple objects can be sent over single TCP connection between client and server
- ❑ HTTP/1.1 uses persistent connections in default mode


Nonpersistent HTTP: Response time

Round Trip Time (RTT:
time to send a small packet to travel from client to server and back.

Response time:

- ❑ one RTT to initiate TCP connection
- ❑ one RTT for HTTP request and first few bytes of HTTP response to return
- ❑ file transmission time

total = $2RTT + \text{transmit time}$



Persistent HTTP

Nonpersistent HTTP issues:

- ❑ requires 2 RTTs per object
- ❑ OS overhead for each TCP connection
- ❑ browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

- ❑ server leaves connection open after sending response
- ❑ subsequent HTTP messages between same client/server sent over open connection

Persistent *without* pipelining:

- ❑ client issues new request only when previous response has been received
- ❑ one RTT for each referenced object

Persistent *with* pipelining:

- ❑ default in HTTP/1.1
- ❑ client sends requests as soon as it encounters a referenced object
- ❑ as little as one RTT for all the referenced objects

HTTP request message

- ❑ two types of HTTP messages: *request, response*
- ❑ **HTTP request message:**
 - ❖ ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

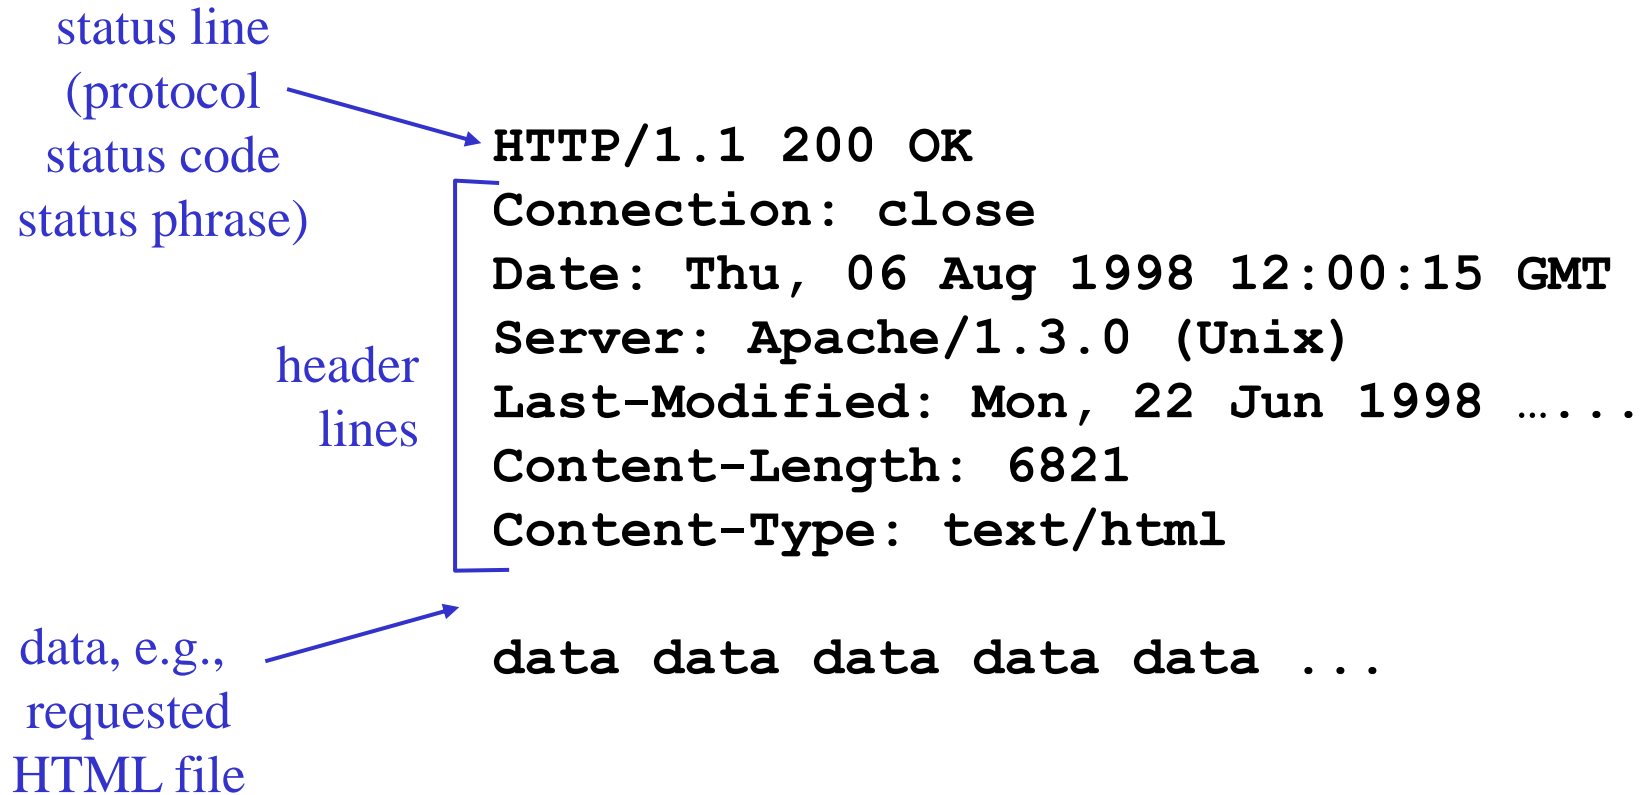
header
lines

Carriage return,
line feed
indicates end
of message

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
(extra carriage return, line feed)
```

Close: nonpersistent connection – 1 TCP
connection per object to send

HTTP response message



HTTP Request - starts with TCP connection

No.	Time	Source	Destination	Protocol	Length	Info
2	0.000374	192.168.50.1	192.168.50.100	TCP	62	58323 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 S
4	0.001036	192.168.50.100	192.168.50.1	TCP	62	80 → 58323 [SYN, ACK] Seq=0 Ack=1 Win=64240 Len=0
5	0.001209	192.168.50.1	192.168.50.100	TCP	54	58323 → 80 [ACK] Seq=1 Ack=1 Win=64240 Len=0
6	0.001851	192.168.50.1	192.168.50.100	HTTP	615	GET / HTTP/1.1
7	0.002524	192.168.50.100	192.168.50.1	TCP	60	80 → 58323 [ACK] Seq=1 Ack=562 Win=63954 Len=0
8	0.003805	192.168.50.100	192.168.50.1	HTTP	473	HTTP/1.1 200 OK (text/html)
9	0.044556	192.168.50.1	192.168.50.100	TCP	54	58323 → 80 [ACK] Seq=562 Ack=420 Win=63821 Len=0

TCP 3-way handshake to initiate the HTTP connect in port 80
 GET: request of HTTP version 1.1
 200 OK: answer with data transmission

Hypertext Transfer Protocol		GET: request of HTTP version 1.1
GET / HTTP/1.1\r\n		
Host: www.arqredes.pt\r\n		The server host with the site
Connection: keep-alive\r\n		Keep-alive: persistent connection to send all objects in the same TCP connection
Cache-Control: max-age=0\r\n		
Upgrade-Insecure-Requests: 1\r\n		
User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/97.0.4692.71 Safari/537.36		
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9		
Sec-GPC: 1\r\n		
Accept-Encoding: gzip, deflate\r\n		
Accept-Language: en-US,en;q=0.9\r\n		
If-None-Match: "4f-5d5ddd3551e77-gzip"\r\n		
If-Modified-Since: Tue, 18 Jan 2022 16:33:13 GMT\r\n		

HTTP Request - answers with data, ACKed with TCP

No.	Time	Source	Destination	Protocol	Length	Info
2	0.000374	192.168.50.1	192.168.50.100	TCP	62	58323 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 S
4	0.001036	192.168.50.100	192.168.50.1	TCP	62	80 → 58323 [SYN, ACK] Seq=0 Ack=1 Win=64240 Len=0
5	0.001209	192.168.50.1	192.168.50.100	TCP	54	58323 → 80 [ACK] Seq=1 Ack=1 Win=64240 Len=0
6	0.001851	192.168.50.1	192.168.50.100	HTTP	615	GET / HTTP/1.1
7	0.002524	192.168.50.100	192.168.50.1	TCP	60	80 → 58323 [ACK] Seq=1 Ack=562 Win=63954 Len=0
8	0.003805	192.168.50.100	192.168.50.1	HTTP	473	HTTP/1.1 200 OK (text/html)
9	0.044556	192.168.50.1	192.168.50.100	TCP	54	58323 → 80 [ACK] Seq=562 Ack=420 Win=63821 Len=0

> HTTP/1.1 200 OK\r\n

Date: Tue, 18 Jan 2022 17:33:07 GMT\r\n

Server: Apache/2.4.38 (Debian)\r\n

Last-Modified: Tue, 18 Jan 2022 16:33:13 GMT\r\n

ETag: "4f-5d5ddd3551e77-gzip"\r\n

Accept-Ranges: bytes\r\n

Vary: Accept-Encoding\r\n

Content-Encoding: gzip\r\n

> Content-Length: 84\r\n

Keep-Alive: timeout=5, max=100\r\n

Connection: Keep-Alive\r\n

Content-Type: text/html\r\n

Keep-alive: persistent connection, all objects are sent in the same TCP connection

TCP 3-way handshake to initiate the HTTP connect in port 80

GET: request of HTTP version 1.1

200 OK: answer sending the data

HTTP: The content of the page

No.	Time	Source	Destination	Protocol	Length	Info
2	0.000374	192.168.50.1	192.168.50.100	TCP	62	58323 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 S
4	0.001036	192.168.50.100	192.168.50.1	TCP	62	80 → 58323 [SYN, ACK] Seq=0 Ack=1 Win=64240 Len=0
5	0.001209	192.168.50.1	192.168.50.100	TCP	54	58323 → 80 [ACK] Seq=1 Ack=1 Win=64240 Len=0
6	0.001851	192.168.50.1	192.168.50.100	HTTP	615	GET / HTTP/1.1
7	0.002524	192.168.50.100	192.168.50.1	TCP	60	80 → 58323 [ACK] Seq=1 Ack=562 Win=63954 Len=0
8	0.003805	192.168.50.100	192.168.50.1	HTTP	473	HTTP/1.1 200 OK (text/html)
9	0.044556	192.168.50.1	192.168.50.100	TCP	54	58323 → 80 [ACK] Seq=562 Ack=420 Win=63821 Len=0

< >

> Frame 8: 473 bytes on wire (3784 bits), 473 bytes captured (3784 bits) on interface \Device\NPF_{0C3DD97D-3D58-4976-978E-B16B8}

> Ethernet II, Src: PcsCompu_8b:31:e1 (08:00:27:8b:31:e1), Dst: 0a:00:27:00:00:03 (0a:00:27:00:00:03)

> Internet Protocol Version 4, Src: 192.168.50.100, Dst: 192.168.50.1

> Transmission Control Protocol, Src Port: 80, Dst Port: 58323, Seq: 1, Ack: 562, Len: 419

> Hypertext Transfer Protocol

√ Line-based text data: text/html (6 lines)

```
<html> \n
\t<body>\n
\t\t<h1>arqredes.pt</h1> \n
\t\t<h2>Porto 80</h2> \n
\t</body> \n
</html>\n
```

arqredes.pt
Porto 80